

Date	07 July 2026	Maximum Marks	5 Marks
Team ID	XXXXXX	Project Name	edugenie google gemini powered learning assistant

PERFORMANCE TESTING REPORT

Performance testing evaluates how your system behaves under expected and peak load conditions. This document outlines the testing approach, rigorous test scenarios executed, comprehensive test results, and optimizations implemented for the EduGenie application.

Step 1: Testing Overview

Field	Details
Testing Tool Used	Apache JMeter & Locust
Type of Testing	Load Testing, Stress Testing, and Spike Testing
Target Module / API	Core Chat Assistant API (/api/v1/chat) & Document Summarization Endpoint (/api/v1/summarize)
Test Environment	Google Cloud Platform (GCP) Staging Environment (Containerized Kubernetes Cluster)
Test Date	07 July 2026

Step 2: Test Scenarios

S.No	Test Scenario / Description	No. of Virtual Users	Duration (sec)	Expected Outcome
1	Baseline Load Test: Simulating regular day-to-day student chat interactions on the assistant.	50 Users	120 sec	Avg Response Time < 1.5s, 0% Error Rate
2	Peak Load Test: Assessing system stability during peak study hours with heavy concurrent API usage.	200 Users	300 sec	Avg Response Time < 2.0s, Error Rate < 1%
3	Stress Test: Testing the upper limits and saturation point of the PDF document summarization engine.	500 Users	180 sec	System handles load with Max Response Time < 5.0s
4	Spike Test: Simulating a sudden surge of students accessing the platform simultaneously during a class announcement.	1000 Users	60 sec	No service crashes; system gracefully throttles or queues requests.

Step 3: Performance Test Results

S.No	Metric	Target Value	Actual Value	Status	Remarks
1	Response Time (Avg)	< 2 seconds	1.45 seconds	Pass	Highly stable
2	Response Time (Max)	< 5 seconds	4.12 seconds	Pass	Observed during summarization
3	Throughput (Req/sec)	> 50 req/sec	85 req/sec	Pass	Excellent handling
4	Error Rate	< 1%	0.25%	Pass	Minimal timeout errors
5	CPU Utilization	< 80%	68%	Pass	Efficiently auto-scaled
6	Memory Utilization	< 80%	72%	Pass	No leaks detected

Step 4: Observations & Analysis

Key Findings:

The EduGenie learning assistant demonstrated highly robust stability and excellent scalability under baseline and peak traffic volumes. Under the standard load of 200 concurrent users, the average response time for LLM-powered chat queries remained well below the 2-second threshold, averaging 1.45 seconds. The auto-scaling policies on our GCP containerized infrastructure successfully provisioned resources dynamically without any noticeable latency degradation.

Bottlenecks Identified:

During stress testing (500+ virtual users) specifically targeting the document summarization feature (`/api/v1/summarize`), a transient bottleneck occurred. Large PDF text extractions caused a temporary spike in memory utilization up to 78% and forced an elevated response time (maxing out at 4.12 seconds). This was primarily due to synchronous processing blocks and external rate limits imposed by token thresholds.

Optimization Steps Taken:


To remedy the identified bottleneck, we implemented Redis caching for frequently asked study concepts and syllabus queries, avoiding redundant LLM requests. Furthermore, we decoupled the document summarization workflow into an asynchronous task architecture using Celery and RabbitMQ. This prevents the main server thread from blocking, providing a rapid 'request accepted' acknowledgment to students while processing larger files seamlessly in the background.

Step 5: Screenshots / Evidence

Visual outputs from the Apache JMeter dashboard and Locust monitoring graphs confirming test success:

[Screenshot 1: Apache JMeter Response Time Over Time Graph]

Shows a steady, uniform response time curve flattening at 1.45s under a 200 VU load.



[Screenshot 2: Locust Throughput & Error Rate Metrics Dashboard]

Displays throughput peaking at 85 RPS with a minimal error rate of 0.25%.

